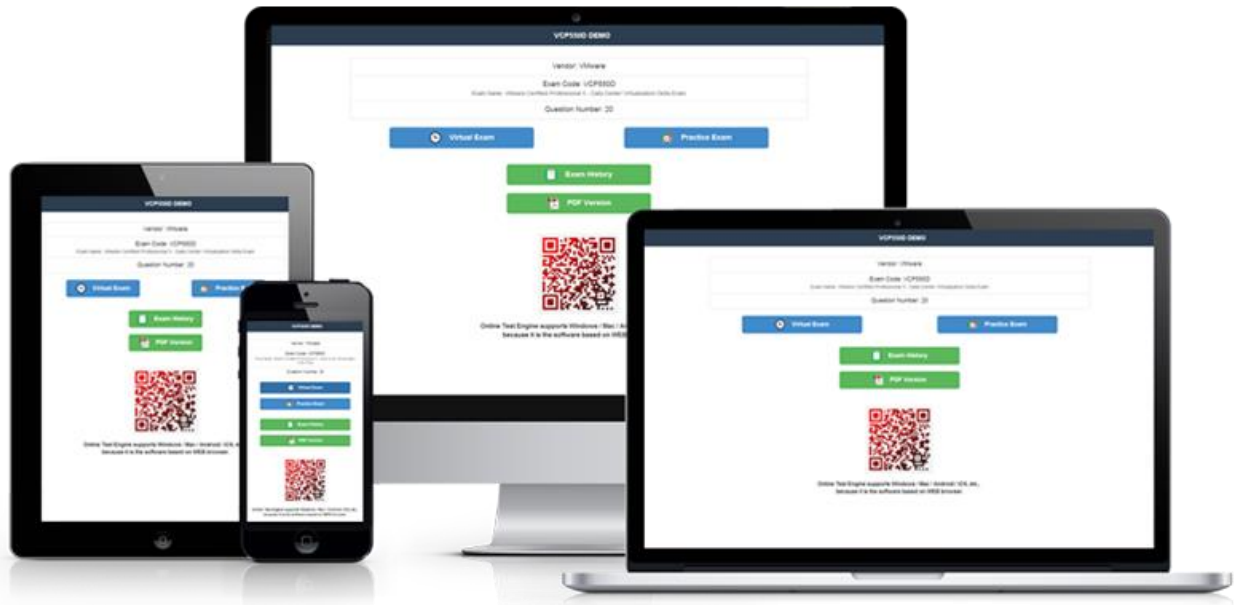


PDFExamDumps



先試後買

購買之前，你可以先嘗試下載一個試用版本。目前我們只提供PDF版本的試用DEMO，軟件版本只提供截圖。

- > 全天候客戶支持，安全的購物網站。
- > 一年免費更新，以符合真正的考試場景。
- > 支付成功以後，你能在網站上立即下載所購買的產品。

<http://www.pdfexamdumps.com>

IT 認證考試題庫-高品質通過率 100%的考試題庫

Exam : **70-767**

Title : Implementing a SQL Data Warehouse

Vendor : Microsoft

Version : DEMO

NO.1 You are developing a Microsoft SQL Server Integration Services (SSIS) package. You enable the SSIS log provider for the Windows event log. You configure the package to use the ScriptTaskLogEntry event. You create a custom Script task.

You need to ensure that when the script completes, it writes the execution status to the event log on the server that hosts SSIS.

Which code segment should you add to the Script task?

- A.** Dts.TaskResult = (int)ScriptResults.Failure
- B.** Dts.Events.FireWarning (0, "SSIS", "Script executed with return result " Dts.TaskResult, String.Empty,0)
- C.** System.Diagnostics.EventLog.writeEntryC'SSIS", "Script executed with return result " + Dts.TaskResult, System.Diagnostics.EventLogEntryType.Information)
- D.** Dts.TaskResult = (int)ScriptResults.Success

Answer: D

NO.2 Note: This question is part of a series of questions that use the same or similar answer choices. An answer choice may be correct for more than one question in the series. Each question is independent of the other questions in this series. Information and details provided in a question apply only to that question.

You are loading data from an OLTP database to a data warehouse. The database contains a table named Sales.

Sales contains details of records that have a type of refund and records that have a type of sales. The data warehouse design contains a table for sales data and a table for refund data.

Which component should you use to load the data to the warehouse?

- A.** the Slowly Changing Dimension transformation
- B.** the Conditional Split transformation
- C.** the Merge transformation
- D.** the Data Conversion transformation
- E.** an Execute SQL task
- F.** the Aggregate transformation
- G.** the Lookup transformation

Answer: B

Explanation

The Conditional Split transformation can route data rows to different outputs depending on the content of the data. The implementation of the Conditional Split transformation is similar to a CASE decision structure in a programming language. The transformation evaluates expressions, and based on the results, directs the data row to the specified output. This transformation also provides a default output, so that if a row matches no expression it is directed to the default output.

References:

<https://docs.microsoft.com/en-us/sql/integration-services/data-flow/transformations/conditionalsplit-Transformation>

NO.3 Note: This question is part of a series of questions that use the same or similar answer choices. An answer choice may be correct for more than one question in the series. Each question is independent of the other questions in this series. Information and details provided in a question

apply only to that question.

You are a database administrator for an e-commerce company that runs an online store. The company has the databases described in the following table.

Database	Description
DB1	This database supports the online store.
DB2	This is the data warehouse for the company. DB2 contains a table named OnlineOrder that is partitioned in hourly increments. The LOCK_ESCALATION option is set to AUTO . The data flow contains 24 OLE DB destinations, one for each partition.
DB3	This database runs Master Data Services (MDS).

Each day, you publish a Microsoft Excel workbook that contains a list of product names and current prices to an external website. Suppliers update pricing information in the workbook. Each supplier saves the workbook with a unique name.

Each night, the Products table is deleted and refreshed from MDS by using a Microsoft SQL Server Integration Services (SSIS) package. All files must be loaded in sequence.

You need to add a data flow in an SSIS package to perform the Excel files import in the data warehouse.


What should you use?

- A. Lookup transformation
- B. Merge transformation
- C. Merge Join transformation
- D. MERGE statement
- E. Union All transformation
- F. Balanced Data Distributor transformation
- G. Sequential container
- H. Foreach Loop container

Answer: A

Explanation

If you're familiar with SSIS and don't want to run the SQL Server Import and Export Wizard, create an SSIS package that uses the Excel Source and the SQL Server Destination in the data flow.



References:

<https://docs.microsoft.com/en-us/sql/integration-services/import-export-data/import-data-from-excel-to-sql>

NO.4 You have a Microsoft SQL Server Integration Services (SSIS) package that loads data into a data warehouse each night from a transactional system. The package also loads data from a set of Comma-Separated Values (CSV) files that are provided by your company's finance department. The SSIS package processes each CSV file in a folder. The package reads the file name for the current file into a variable and uses that value to write a log entry to a database table. You need to debug the package and determine the value of the variable before each file is processed. Which four actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions

Click the **Start** toolbar button to commence debugging the package.

When a breakpoint is reached, view the value of the variable by using the Variables window.

Open the Control Flow editor for the package.

When a breakpoint is reached, view the value of the variable by using the Locals window.

Set a breakpoint on the For Loop container.

Set a breakpoint on the Sequence container.

Open the Data Flow editor for the package.

Set a breakpoint on the Foreach Loop container.

Answer Area



Answer:

Actions

Click the **Start** toolbar button to commerce debugging the package.

When a breakpoint is reached, view the value of the variable by using the Variables window.

Open the Control Flow editor for the package.

When a breakpoint is reached, view the value of the variable by using the Locals window.

Set a breakpoint on the For Loop container.

Set a breakpoint on the Sequence container.

Open the Data Flow editor for the package.

Set a breakpoint on the Foreach Loop container.

Answer Area

Open the Control Flow editor for the package.

Set a breakpoint on the Foreach Loop container.



Click the **Start** toolbar button to commerce debugging the package.



When a breakpoint is reached, view the value of the variable by using the Locals window.



Explanation

Answer Area

Open the Control Flow editor for the package.

Set a breakpoint on the Foreach Loop container.

Click the **Start** toolbar button to commerce debugging the package.

When a breakpoint is reached, view the value of the variable by using the Locals window.

You debug control flows.

The Foreach Loop container is used for looping through a group of files. Put the breakpoint on it. The Locals window displays information about the local expressions in the current scope of the

Transact-SQL debugger.

References: <https://docs.microsoft.com/en-us/sql/integration-services/troubleshooting/debugging-control-flow>

<http://blog.pragmaticworks.com/looping-through-a-result-set-with-the-foreach-loop>

NO.5 Note: This question is part of a series of questions that use the same or similar answer choices. An answer choice may be correct for more than one question in the series. Each question is independent of the other questions in this series. Information and details provided in a question apply only to that question.

You are designing a data warehouse and the load process for the data warehouse.

You have a source system that contains two tables named Table1 and Table2. All the rows in each table have a corresponding row in the other table.

The primary key for Table1 is named Key1. The primary key for Table2 is named Key2.

You need to combine both tables into a single table named Table3 in the data warehouse. The solution must ensure that all the nonkey columns in Table1 and Table2 exist in Table3.

Which component should you use to load the data to the data warehouse?

- A. the Slowly Changing Dimension transformation
- B. the Conditional Split transformation
- C. the Merge transformation
- D. the Data Conversion transformation
- E. an Execute SQL task
- F. the Aggregate transformation
- G. the Lookup transformation

Answer: G

Explanation

The Lookup transformation performs lookups by joining data in input columns with columns in a reference dataset. You use the lookup to access additional information in a related table that is based on values in common columns.

You can configure the Lookup transformation in the following ways:

Specify joins between the input and the reference dataset.

Add columns from the reference dataset to the Lookup transformation output.

Etc.

NO.6 You need to load data from a CSV file to a table.

How should you complete the Transact-SQL statement? To answer, drag the appropriate Transact-SQL segments to the correct locations. Each Transact-SQL segment may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Transact-SQL segments

BULK	INSERT
FROM	WITH
MERGE	

Answer Area

		Sales.Invoices
		'\\share\data\file1.csv'
		(FORMAT = 'CSV')

Answer:**Transact-SQL segments**

BULK	INSERT
FROM	WITH
MERGE	

Answer Area

BULK	INSERT	Sales.Invoices
FROM		'\\share\data\file1.csv'
WITH		(FORMAT = 'CSV')

Explanation

Answer Area

BULK	INSERT	Sales.Invoices
FROM		'\\share\data\file1.csv'
WITH		(FORMAT = 'CSV')

Example:

```
BULK INSERT Sales.Orders
FROM '\\SystemX\DiskZ\Sales\data\orders.csv'
WITH ( FORMAT='CSV');
```

NO.7 Your company manufactures several types of products.

The company has a production tracking application that stores the following data about the products:

The production date

The cost of production

The names of the products

The amount of waste created

The number of products produced

The name of the facility where the products are produced

You are designing a data warehouse for the data. You add a Date dimension.

You need to ensure that you can create a composite primary key for the fact table.

Which two columns should you add to the new dimension tables? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Cost of Production
- B. Amount Produced
- C. Waste Amount
- D. Product Name
- E. Facility Name

Answer: D E

NO.8 You have a data warehouse that includes a table named [Fact].[sales]. The table has four partitions. There are no foreign keys defined on the [Fact].[sales] table.

You identify an issue with the first partition. You must reload the data for the first partition as quickly as possible.

You need to prepare the partition for reloading data.

How should you complete the Transact-SQL statement? To answer, drag the appropriate Transact-SQL segments to the correct locations. Each Transact-SQL segment may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

Answer:

NO.9 You are developing a Microsoft SQL Server Master Data Services (MDS) solution.

The model contains an entity named Product. The Product entity has three user-defined attributes named Category, Subcategory, and Price, respectively.

You need to ensure that combinations of values stored in the Category and Subcategory attributes are unique.

What should you do?

- A. Create an attribute group that consists of the Category and Subcategory attributes. Publish a

business rule for the attribute group.

B. Publish a business rule that will be used by the Product entity.

C. Create a derived hierarchy based on the Category and Subcategory attributes. Use the Category attribute as the top level for the hierarchy.

D. Set the value of the Attribute Type property for the Category and Subcategory attributes to Domainbased.

Answer: B

Explanation

In Master Data Services, business rule actions are the consequence of business rule condition evaluations. If a condition is true, the action is initiated.

The Validation action "must be unique": The selected attribute must be unique independently or in combination with defined attributes.

NO.10 Note: This question is part of a series of questions that use the same or similar answer choices. An answer choice may be correct for more than one question in the series. Each question is independent of the other questions in this series. Information and details provided in a question apply only to that question.

You are developing a Microsoft SQL Server Integration Services (SSIS) package.

You need to ensure that the packa

ge records the current Log Sequence Number (LSN) in the source database before the package begins reading source tables.

Which SSIS Toolbox item should you use?

A. CDC Control task

B. CDC Splitter

C. Union All

D. XML task

E. Fuzzy Grouping

F. Merge

G. Merge Join

Answer: A

Explanation

The CDC Control task is used to control the life cycle of change data capture (CDC) packages. It handles CDC package synchronization with the initial load package, the management of Log Sequence Number (LSN) ranges that are processed in a run of a CDC package.

References: <https://docs.microsoft.com/en-us/sql/integration-services/control-flow/cdc-control-task>

NO.11 You manage a data warehouse. You have a source table named dbo.Table1.

You need to implement Change Data Capture (CDC) for the queries.

Query name	Requirement
QueryA	Return the primary key values for any rows that were updated or inserted between two datetime values.
QueryB	Return all update and insert operations, including the before and after values for any update statements, between two datetime values.
QueryC	Review all update and insert operations. Only the new values for update operations between two datetime values are returned.

Which CDC object should you use for each query? To answer, drag the appropriate objects to the correct queries. Each object may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Answer Area

Objects

dbo.fn_all_changes_Table1
dbo.fn_net_changes_Table1
sys.fn_cdc_map_lsn_to_time
sys.fn_cdc_map_time_to_lsn

Query name

Object

QueryA

QueryB

QueryC

Answer:

Answer Area

Objects

dbo.fn_all_changes_Table1
dbo.fn_net_changes_Table1
sys.fn_cdc_map_lsn_to_time
sys.fn_cdc_map_time_to_lsn

Query name

Object

QueryA

QueryB

QueryC

dbo.fn_net_changes_Table1

dbo.fn_all_changes_Table1

dbo.fn_all_changes_Table1

Explanation

Query name**Object**

QueryA

dbo.fn_net_changes_Table1

QueryB

dbo.fn_all_changes_Table1

QueryC

dbo.fn_all_changes_Table1

References:

<https://docs.microsoft.com/en-us/sql/relational-databases/system-functions/sys-fn-net-changes-capture-instance-t>

<https://docs.microsoft.com/en-us/sql/relational-databases/system-functions/sys-fn-all-changes-capture-instance-tr>

<https://docs.microsoft.com/en-us/sql/relational-databases/system-functions/sys-fn-cdc-map-lsn-to-time-transact-s>

<https://docs.microsoft.com/en-us/sql/relational-databases/system-functions/cdc-fn-cdc-get-net-changes-capture-in>

NO.12 Note: This question is part of a series of questions that use the same or similar answer choices. An answer choice may be correct for more than one question in the series. Each question is independent of the other questions in this series. Information and details provided in a question apply only to that question.

You are developing a Microsoft SQL Server Integration Services (SSIS) package. The package design consists of the sources shown in the following diagram:



Each source contains data that is not sorted.

You need to combine data from all of the sources into a single dataset.

Which SSIS Toolbox item should you use?

- A. CDC Control task
- B. CDC Splitter
- C. Union All
- D. XML task
- E. Fuzzy Grouping

F. Merge

G. Merge Join

Answer: C